

A Pilot Stress Test of LLM-as-a-Judge Robustness in Culturally Sensitive Annotation

Maksym Yemelianenko* Wesley Zhuang
NYU Tandon School of Engineering NYU Tandon School of Engineering
Dorien Zhang
NYU Tandon School of Engineering

June 2026

Abstract

Large language models (LLMs) are increasingly used as automated judges for benchmarking, preference ranking, and large-scale annotation. Their reliability is less clear in subjective and culturally sensitive tasks, where human disagreement may reflect meaningful variation rather than annotation noise. We present a pilot diagnostic stress test of LLM-as-a-Judge robustness for offensiveness annotation. The goal is not to estimate population-level flip rates, but to test whether subtle cultural and rhetorical framing can induce measurable judgment instability under controlled paired comparisons.

We evaluate five LLM judges available at the time of evaluation—Claude Haiku 4.5, GPT-5.4, Gemini 3.5 Flash, Llama 3.1 8B, and Qwen3 32B—on a deliberately small, balanced subset of 30 English examples from the SemEval-2023 Learning with Disagreements offensiveness data. Each item is evaluated under one control condition and nine framing perturbations, with three repeated calls per item-condition pair. The main analysis aggregates the three calls by majority vote, producing 270 paired perturbation comparisons per model; raw per-call analyses are reported as robustness checks.

The pilot reveals possible instability across all models. Majority-vote label-flip rates range from 6.7% for Gemini 3.5 Flash to 17.0% for Claude Haiku 4.5, with Wilson 95% intervals that remain wide because of the small sample. Instability is highest for examples with mixed human judgments: items with 40–60% human offensive votes show majority-vote flip rates around 21–22%, compared with 2.2% for unanimously offensive items. A four-reviewer semantic validation study finds that most perturbations preserve meaning reasonably well, while west-comparison, in-group pride, west-bias, and traditionalism wrappers are less semantically invariant and should be interpreted as framing interventions rather than strictly meaning-preserving edits. A neutral-wrapper baseline produces lower majority-vote flip rates than the full perturbation suite for every model, though intervals remain wide. After Benjamini–Hochberg correction across model–perturbation comparisons, only one individual comparison remains significant: Claude Haiku 4.5 under west-comparison framing. These results should therefore be read as preliminary evidence of possible fragility, not as a definitive benchmark of model fairness or moderation reliability. They motivate larger-scale follow-up studies with broader multilingual coverage.

Keywords: LLM-as-a-Judge; pilot study; stress testing; human disagreement; cultural sensitivity; offensiveness detection; robustness.

Artifacts: <https://github.com/1mpossible-code/llm-judge-stress-test>

*Corresponding author.

1 Introduction

Large language models (LLMs) are increasingly used not only to generate text, but also to evaluate it. LLM-as-a-Judge systems now appear in model leaderboards, preference ranking pipelines, reward-model construction, red-teaming, and annotation workflows. Their appeal is clear: compared with human evaluation, they are fast, scalable, and easy to standardize. Prior work has shown that strong LLMs can correlate with human preferences in some open-ended evaluation settings [4]. As a result, LLM judges are becoming part of the infrastructure through which models are selected, aligned, and deployed.

However, giving models evaluative authority raises a reliability question: are LLM judges stable in domains where human judgment is subjective, culturally contingent, and legitimately contested? Tasks such as offensiveness detection, hate speech classification, sexism detection, political norm interpretation, and cultural sensitivity annotation often lack a single unambiguous ground truth. In these settings, disagreement among annotators may reflect differences in background, identity, pragmatic inference, social norms, or tolerance for harmful speech [10, 11, 13]. Treating such disagreement as noise can erase minority interpretations or create false consensus.

This paper reframes the problem as a pilot robustness stress test. Because this study uses a deliberately small, balanced subset of 30 examples, the results should be interpreted as a diagnostic stress test rather than a population-level estimate of model behavior. Our goal is to examine whether culturally framed perturbations can induce instability in LLM-as-Judge offensiveness decisions under controlled conditions. We therefore emphasize paired within-item comparisons, uncertainty intervals, self-consistency checks, and cautious interpretation.

We ask three research questions:

- RQ1: Under a majority-vote repeated-call protocol, do LLM judges show label instability when the same base text is presented with cultural or rhetorical framing?
- RQ2: Are observed flips concentrated in particular perturbation types or human-disagreement levels?
- RQ3: How much of the observed instability appears to exceed baseline repeated-call variation within the same condition?

1.1 Contributions

This work makes four contributions:

- We introduce a small paired stress-test protocol for evaluating LLM-as-a-Judge robustness in a culturally sensitive annotation task.
- We report uncertainty estimates, majority-vote analyses, raw repeated-call robustness checks, self-consistency measures, paired McNemar tests, and false-discovery-rate adjusted p-values.
- We analyze instability by model, perturbation type, human agreement level, and individual item, showing that the strongest pilot effects occur in high-disagreement examples and a small number of model-perturbation cells.
- We release reproducible artifacts including selected item IDs, prompt templates, perturbation templates, raw outputs, parsed judgments, paired comparisons, and analysis scripts.

2 Related Work

2.1 LLM-as-a-Judge

LLM-based evaluation has become a prominent alternative to human evaluation for open-ended generation tasks. Early work showed that LLMs can serve as reference-free evaluators for generation and translation quality [1, 2]. Zheng et al. [4] introduced MT-Bench and Chatbot Arena, showing that strong LLM judges can approximate human preferences while also identifying limitations such as position and verbosity bias. Recent surveys characterize LLM-as-a-Judge as a rapidly expanding evaluation paradigm with unresolved reliability and bias challenges [3]. Follow-up work has studied positional sensitivity [5, 8], verbosity bias [6], self-preference [9, 7], and broader evaluation fragility [7]. This literature motivates our premise: if LLMs are used as evaluators, the evaluators themselves need stress testing.

2.2 Human Disagreement and Subjective Annotation

Human label variation is increasingly recognized as a central issue in NLP evaluation [13]. In subjective tasks, annotator disagreement can reflect valid interpretive diversity rather than annotator error. Pavlick and Kwiatkowski [12], for example, show that some inference disagreements persist even after additional judgments and context variation. Work on learning from disagreement argues that preserving multiple annotations can improve modeling and evaluation for socially grounded tasks [10, 11, 14]. The SemEval Learning with Disagreements tasks provide benchmark datasets explicitly designed around this perspective [10, 11].

2.3 Cultural Framing and Moderation Sensitivity

Offensiveness and harmful-speech judgments are culturally and pragmatically sensitive. Prior work on hate-speech detection shows that annotator and model judgments can encode social and dialectal bias; Sap et al. [15] show that dialect markers can be spuriously associated with toxicity labels and downstream classifier behavior. The perceived severity of a text may depend on speaker identity, audience, local norms, political context, rhetorical genre, and social power. Our perturbations are proxies for cultural and rhetorical framing, not claims about real cultural identity. We use them to test whether LLM judgments shift when such cues are introduced.

3 Experimental Design

3.1 Dataset and Scope

We use the offensiveness detection data from SemEval-2023 Task 11: Learning with Disagreements (LeWiDi) [11]. Each example has five human annotations, allowing us to represent disagreement as an empirical distribution. We construct a deterministic 30-example English subset balanced by human offensive-label rate: five examples each at $\{0.0, 0.2, 0.4, 0.6, 0.8, 1.0\}$. This sampling intentionally over-represents both consensus and high-disagreement cases. It is useful for diagnosis, but it is not representative of the full dataset distribution.

The sample size was chosen for pilot scope and cost control. The experiment requires repeated calls across five models, ten main-suite conditions, neutral-baseline conditions, and repair/retry passes for provider failures. A 30-item design keeps the study auditable by humans while still covering all six human-disagreement strata. We therefore use the sample to identify whether framing sensitivity appears at all, not to estimate prevalence in the full LeWiDi population.

3.2 Units of Analysis

We use the following terminology throughout the paper:

- Base item: one original text from LeWiDi.
- Condition: the control text or one perturbed version of the same item.
- Trial: one model response to one item under one condition.
- Repetition: repeated sampling of the same item-condition pair. We use three repetitions.
- Majority label: the majority vote across the three repeated labels for one model, item, and condition.
- Flip: a perturbation majority label differs from the matched control majority label for the same model and base item.

The main results use majority labels because they reduce the influence of within-condition stochasticity. The robustness appendix also reports raw per-trial flip rates, where each perturbation trial is paired with the corresponding control trial of the same repetition index.

3.3 Perturbation Taxonomy

We evaluate nine perturbation conditions plus the unmodified control. Table 1 summarizes the intended mechanism for each perturbation family. The taxonomy is diagnostic rather than exhaustive; perturbations are designed to probe plausible routes by which framing may alter an LLM judge’s interpretation.

Perturbation family	Mechanism probed	Possible effect
West comparison	Cross-cultural norm anchoring and comparison framing	May increase or decrease strictness depending on norm interpretation
In-group pride	Identity and community-affirmation cue	May soften perceived harm or increase defensiveness
Cultural context	Explicit claim that the statement is understood differently in a local context	May shift threshold for offensiveness
Traditionalism	Cultural preservation / modern-decay framing	May sanitize or intensify moral judgment
Authority	Expert-consensus cue	May increase deference to the wrapper framing
Verbosity	Known LLM-judge sensitivity to fluent padding	May alter salience independent of base content
Position/polarity	Affirmative or negative lead-in framing	May probe anchoring and polarity effects

Table 1: Perturbation taxonomy and intended diagnostic mechanisms.

3.4 Human Validation of Perturbations

To assess whether perturbations preserve the original meaning, we conducted a small human validation study. Four reviewers independently rated all 270 original–perturbed pairs on three 1–5 scales: meaning preservation, whether the edit added only framing, and whether the edit directly changed perceived offensiveness. The interface randomized item order, hid perturbation metadata by default, and repeated 10% of pairs as blinded duplicate reliability checks. Duplicate-check mean absolute differences were 0.56 for meaning preservation, 0.54 for framing-only ratings, and 0.57 for direct-offensiveness-change ratings; 86–91% of repeated ratings were within one scale point.

Table 2 reports validation results. Position framing, verbosity, authority, and in-group cultural-context wrappers score above 4.0 on meaning preservation. West-comparison, in-group pride, west-bias, and traditionalism wrappers score below our 4.0 preservation analysis threshold, although their direct-offensiveness-change scores remain low. We therefore interpret these perturbations as controlled framing interventions rather than strictly semantic-preserving edits.

Factor	Level	n	Meaning	Framing-only	Direct offense change
authority	expert_consensus	120	4.22	3.88	1.52
cultural_framing	ingroup_medium	120	4.12	4.01	1.42
cultural_framing	west_bias_medium	120	3.96	3.62	1.56 [†]
cultural_pride	ingroup_pride	120	3.95	3.78	1.44 [†]
cultural_pride	west_comparison	120	3.77	3.57	1.82 [†]
position_bias	affirmative_medium	120	4.67	4.78	1.09
position_bias	negative_medium	120	4.11	4.30	1.38
traditionalism	modern_decay	120	3.76	3.56	1.63 [†]
verbosity	padded_fluent	120	4.43	4.52	1.27

Table 2: Human semantic validation of perturbations. Scores are means on 1–5 scales across four reviewers and 30 examples per perturbation ($n = 120$ ratings each). Higher meaning and framing-only scores are better; lower direct-offensiveness-change scores are better. [†] indicates mean meaning preservation below 4.0 or direct-offensiveness-change above 2.5.

3.5 Models

We evaluate five LLM judges available at the time of evaluation. Table 3 lists the model identifiers reported for the experiment. Interface-specific access details are treated as implementation metadata and are not used as model names. Model references are provided for reproducibility [21, 22, 23, 19, 20].

Display name	Reported model identifier
Claude Haiku 4.5	claude-haiku-4-5-20251001
GPT-5.4	gpt-5.4 / snapshot gpt-5.4-2026-03-05
Gemini 3.5 Flash	gemini-3.5-flash
Llama 3.1 8B	llama-3.1-8b-instant
Qwen3 32B	qwen/qwen3-32b

Table 3: Judge models and reported identifiers used in the experiments.

3.6 Protocol and Randomness

For each model, we evaluate 30 base items under 10 main-suite conditions: one control and nine perturbations. Each item-condition pair is repeated three times, yielding 900 judgments per model

and 4,500 main-suite judgments total. The main paired analysis compares $30 \times 9 = 270$ majority-label perturbation pairs per model. Raw per-trial robustness checks compare $30 \times 9 \times 3 = 810$ paired trials per model. We additionally run a neutral-wrapper baseline with one control and one neutral condition, adding 180 judgments per model and 900 judgments overall.

The perturbation selection seed is 1337. The judge prompt is fixed across models, with prompt hash 6b67535b23a51992a5e09ca814fde4aa9bf3a12935a23e09d986c85c1b33b403. Temperature is set to 0.0 for all runs. Max output tokens are 64 for four models and 160 for Qwen3 32B because its responses required more room for parseable rationales. API-level random seeds and top-p controls were not available uniformly across all providers, so repeated calls are treated as samples from each model’s operational judgment distribution. Calls were made on June 25–26, 2026. Raw outputs are cached; failed calls are rerun for the same item-condition cell, while successful judgments are not manually edited.

4 Metrics and Statistical Analysis

4.1 Label Instability

For a model m , item i , perturbation p , and majority label function $Y(\cdot)$, a flip occurs when

$$Y_m(x_i^{(p)}) \neq Y_m(x_i^{(0)}), \quad (1)$$

where $x_i^{(0)}$ is the control condition. Attack Success Rate (ASR), used here simply as a label-flip rate, is

$$\text{ASR} = \frac{\#\text{perturbation comparisons where label differs from control}}{\#\text{total perturbation comparisons}}. \quad (2)$$

We report Wilson 95% confidence intervals for all ASR estimates [16]. Because the sample is small and balanced by design, these intervals should be interpreted as uncertainty for the pilot sample, not as population confidence intervals for the full LeWiDi distribution.

4.2 Distributional Divergence

Human annotations are represented as an empirical distribution H_i over {offensive, not offensive}. For a model majority label, we use a one-hot distribution M_i . Jensen–Shannon divergence is computed as

$$\text{JSD}(M_i, H_i) = \frac{1}{2}\text{KL}(M_i||A_i) + \frac{1}{2}\text{KL}(H_i||A_i), \quad A_i = \frac{M_i + H_i}{2}. \quad (3)$$

With base-2 logarithms, JSD ranges from 0 to 1. Larger values indicate greater divergence between the model’s categorical judgment and the human label distribution. We report

$$\Delta\text{JSD}_p = \text{JSD}(M_i^{(p)}, H_i) - \text{JSD}(M_i^{(0)}, H_i). \quad (4)$$

Positive ΔJSD means the perturbation moves the model farther from the human distribution; negative values mean it moves closer. We interpret ΔJSD comparatively rather than as an absolute harm threshold.

4.3 Paired Tests and Multiple Comparisons

Because each perturbation is paired with a control judgment for the same base item, we use paired tests rather than treating observations as independent. For each model–perturbation pair, we apply McNemar’s exact test to the majority-label control and perturbation outcomes [17]. We report raw p-values and Benjamini–Hochberg false-discovery-rate adjusted q-values across the 45 model–perturbation comparisons [18]. Since our analysis compares multiple models and perturbation types, we treat uncorrected comparisons as exploratory. For confidence changes, we use Wilcoxon signed-rank tests over paired confidence differences and again report BH-adjusted values.

5 Results

5.1 Self-Consistency Across Repeated Calls

Before attributing changes to perturbations, we estimate baseline repeated-call stability. Table 4 reports the proportion of item-condition cells in which all three repeated labels agree. Control-only unanimity ranges from 86.7% for Claude Haiku 4.5 to 100% for Llama 3.1 8B. This indicates that most cells are internally stable, but some observed flips may still reflect operational stochasticity rather than perturbation-specific effects.

Model	Conditions	Unanimity	Control unanimity	Pairwise agreement
Claude Haiku 4.5	300	87.3%	86.7%	91.6%
GPT-5.4	300	84.3%	90.0%	89.6%
Gemini 3.5 Flash	300	86.3%	93.3%	90.9%
Llama 3.1 8B	300	99.7%	100.0%	99.8%
Qwen3 32B	300	98.7%	96.7%	99.1%

Table 4: Repeated-call self-consistency by model. Unanimity is the share of item-condition cells where all three repeated labels match.

5.2 Label Instability by Model

Table 5 reports majority-vote ASR by model. The highest observed pilot flip rates are Claude Haiku 4.5 at 17.0% and Llama 3.1 8B at 16.7%. GPT-5.4 and Gemini 3.5 Flash show lower majority-vote flip rates, 7.8% and 6.7%, respectively. The intervals are intentionally reported and remain wide, reflecting the small diagnostic sample.

Model	n	Flips	ASR	95% CI	Direction
Claude Haiku 4.5	270	46	17.0%	[13.0, 22.0]	80.4% not→off
GPT-5.4	270	21	7.8%	[5.1, 11.6]	71.4% not→off
Gemini 3.5 Flash	270	18	6.7%	[4.3, 10.3]	66.7% not→off
Llama 3.1 8B	270	45	16.7%	[12.7, 21.6]	84.4% off→not
Qwen3 32B	270	34	12.6%	[9.2, 17.1]	58.8% off→not

Table 5: Majority-vote label instability by model. The unit is one model–item–perturbation pair after aggregating three repeated calls per condition.




Figure 1: Majority-vote ASR by model with Wilson 95% intervals.

Raw per-trial results are similar in ordering but slightly higher for GPT-5.4 and Gemini 3.5 Flash: Claude Haiku 4.5 17.5%, Llama 3.1 8B 16.8%, Qwen3 32B 11.9%, GPT-5.4 10.1%, and Gemini 3.5 Flash 9.5%. We therefore use majority-vote results as the main analysis and raw per-call results as a robustness check.

5.3 Neutral Wrapper Baseline

To test whether models flip merely because extra text is added, we ran an additional neutral-wrapper baseline. The neutral condition prepends a source/context statement—“The following text is from an online discussion:”—without adding cultural, authority, traditionalist, or polarity cues. Table 6 compares majority-vote neutral ASR with the full perturbation-suite ASR from Table 5. Neutral ASR is lower than the full perturbation-suite ASR for every model, suggesting that the observed effects are not explained solely by adding wrapper text. However, the neutral baseline also produces nonzero flips for four models and confidence intervals overlap for several comparisons, so this result should be treated as supportive rather than conclusive.

Model	n	Neutral flips	Neutral ASR	95% CI	Core ASR
Claude Haiku 4.5	30	4	13.3%	[5.3, 29.7]	17.0%
GPT-5.4	30	2	6.7%	[1.8, 21.3]	7.8%
Gemini 3.5 Flash	30	1	3.3%	[0.6, 16.7]	6.7%
Llama 3.1 8B	30	0	0.0%	[0.0, 11.4]	16.7%
Qwen3 32B	30	2	6.7%	[1.8, 21.3]	12.6%

Table 6: Neutral-wrapper baseline compared with the full perturbation suite. The neutral condition adds source/context text without cultural or evaluative framing.

5.4 Perturbation Effects and Multiple Comparisons

Across all models, west-comparison framing has the highest majority-vote ASR at 16.7%, followed by in-group pride at 14.0%, west-bias framing and authority framing at 13.3%, and traditionalism at 12.7% (Table 7). These differences are suggestive rather than definitive; intervals overlap substantially. The human validation results in Table 2 are important for interpretation: several

high-ASR cultural/traditionalism wrappers were also rated as less meaning-preserving, so we treat them as culturally loaded framing interventions rather than pure invariance tests.

Perturbation	n	Flips	ASR	95% CI
West comparison	150	25	16.7%	[11.6, 23.4]
In-group pride	150	21	14.0%	[9.3, 20.5]
West-bias framing	150	20	13.3%	[8.8, 19.7]
Authority	150	20	13.3%	[8.8, 19.7]
Traditionalism	150	19	12.7%	[8.3, 18.9]
Verbosity	150	18	12.0%	[7.7, 18.2]
Affirmative framing	150	18	12.0%	[7.7, 18.2]
In-group framing	150	12	8.0%	[4.6, 13.5]
Negative framing	150	11	7.3%	[4.1, 12.7]

Table 7: Majority-vote ASR by perturbation type, aggregated across models. The unit is one model-item pair for a given perturbation.




Figure 2: Majority-vote ASR heatmap by model and perturbation. Cell values are percentages rounded to the nearest integer.

Table 8 lists the ten largest model-perturbation ASR cells and their paired-test results. After BH correction across 45 tests, only Claude Haiku 4.5 under west-comparison framing remains below $q < 0.05$. This does not prove a general cultural bias mechanism; it indicates that, in this pilot, one model-framing combination produced a statistically distinguishable directional shift under paired testing.

Model	Perturbation	n	ASR	95% CI	p	BH q
Claude Haiku 4.5	West comparison	30	36.7%	[21.9, 54.5]	0.001	0.044
Claude Haiku 4.5	In-group pride	30	26.7%	[14.2, 44.4]	0.070	0.527
Llama 3.1 8B	Traditionalism	30	23.3%	[11.8, 40.9]	0.016	0.352
Llama 3.1 8B	In-group pride	30	23.3%	[11.8, 40.9]	0.125	0.562
Llama 3.1 8B	Verbosity	30	20.0%	[9.5, 37.3]	0.031	0.469
Claude Haiku 4.5	Authority	30	20.0%	[9.5, 37.3]	0.219	0.865
Llama 3.1 8B	West-bias framing	30	16.7%	[7.3, 33.6]	0.062	0.527
Qwen3 32B	West comparison	30	16.7%	[7.3, 33.6]	0.062	0.527
Claude Haiku 4.5	West-bias framing	30	16.7%	[7.3, 33.6]	1.000	1.000
Llama 3.1 8B	Authority	30	16.7%	[7.3, 33.6]	1.000	1.000

Table 8: Largest model-perturbation ASR cells with McNemar exact p-values and Benjamini-Hochberg adjusted q-values. Uncorrected results are exploratory.

5.5 Instability by Human Agreement Level

The balanced sample lets us examine whether instability concentrates in ambiguous examples. Table 9 and Figure 3 show that examples with mixed human labels have the highest observed ASR: 20.9% for 40% offensive items and 22.2% for 60% offensive items. In contrast, unanimously offensive items have a 2.2% flip rate. This pattern suggests that LLM judges may be especially framing-sensitive near the human decision boundary, although the pilot sample is too small to claim a general monotonic relationship.

Human offensive rate	n	Flips	ASR	95% CI
0%	225	20	8.9%	[5.8, 13.3]
20%	225	15	6.7%	[4.1, 10.7]
40%	225	47	20.9%	[16.1, 26.7]
60%	225	50	22.2%	[17.3, 28.1]
80%	225	27	12.0%	[8.4, 16.9]
100%	225	5	2.2%	[1.0, 5.1]

Table 9: Majority-vote ASR by human offensive-label rate. Each level contains five base items, evaluated across five models and nine perturbations.




Figure 3: Majority-vote ASR by human offensive-label rate. Shaded intervals are Wilson 95% intervals.

Per-item analysis also shows concentration: the most unstable item accounts for 16 flips out of 45 possible model-perturbation comparisons, while several items show no flips. A full item-level table with IDs, human vote counts, flip counts, and most frequent flip perturbation is included in Appendix C. Leave-one-out analysis indicates that headline model ASRs vary by at most 3.1 percentage points when any single item is removed, suggesting that the overall ranking is not entirely driven by one example.

5.6 Confidence and Rationale Instability

Confidence changes are smaller than label changes but still informative. Table 10 reports mean signed and absolute confidence shifts after majority aggregation. Four models show small negative signed shifts with BH-adjusted Wilcoxon values below 0.05, while Gemini 3.5 Flash shows no clear signed shift. These confidence changes should be interpreted cautiously: confidence scales are model-specific and not calibrated across providers.

Model	n	Mean Δc	Mean $ \Delta c $	BH-adjusted p
Claude Haiku 4.5	270	-0.012	0.050	0.026
GPT-5.4	270	-0.010	0.018	0.000
Gemini 3.5 Flash	270	0.001	0.019	0.887
Llama 3.1 8B	270	-0.016	0.019	0.000
Qwen3 32B	270	-0.015	0.029	0.000

Table 10: Confidence instability by model. Δc is perturbed mean confidence minus control mean confidence after aggregating three repeated calls.

Qualitatively, rationales often changed even when labels did not. In several cases, cultural-context wrappers caused models to emphasize contextual interpretation rather than surface insult severity. In other cases, fluent padding or traditionalist framing led the model to describe direct insults as argumentative rather than abusive. These examples indicate possible rationale drift, but we do not claim a formal taxonomy of rationale changes without additional human coding.

5.7 Qualitative Case Studies

Table 11 gives three concrete examples of label and rationale shifts. We summarize or redact offensive text to reduce exposure while preserving the diagnostic pattern. The examples are illustrative rather than exhaustive; they are selected from high- Δ JSD and high-confidence-swing exhibits.

Case	Control text summary	Perturbation	Label / confidence shift	Rationale shift
Political accusation	A political post accuses a public figure of inciting violence and uses a punitive hashtag.	West-comparison framing: “Western perspectives often miss that ...”	GPT-5.4: not offensive / 0.87 $\rightarrow$ offensive / 0.91	The control rationale treats the text as political criticism; the perturbed rationale emphasizes hostile and inflammatory tone.
Religious insult	A post uses religious language to tell a named person to “shut up” and go “to hell.”	In-group pride/context framing: “Our community has always understood that ...”	Gemini 3.5 Flash: offensive / 0.85 $\rightarrow$ not offensive / 0.85	The perturbed rationale frames the wording as culturally contextualized and below the threshold for severe abuse.
Direct insult in argument	A post contains direct insults, including a redacted vulgar insult, while making an argumentative political/legal point.	Verbose framing: “It is important to note that ...”	Llama 3.1 8B: offensive / 0.80 $\rightarrow$ not offensive / 0.80	The perturbed rationale gives more weight to argument structure and treats the insult as provocative rather than abusive.

Table 11: Illustrative qualitative cases. Text is summarized or partially redacted; labels and confidences are from paired control–perturbation exhibits.

6 Discussion

This pilot provides preliminary evidence that LLM judges can be sensitive to cultural and rhetorical framing in a subjective annotation task. The evidence is strongest as a diagnostic signal: all models show some majority-vote label flips, flips are more frequent in human-disagreement regions, and one model–perturbation comparison survives FDR correction. The human validation study strengthens the design by identifying which perturbations are closer to meaning-preserving wrappers and which are better understood as stronger framing interventions. At the same time, the study does not establish population-level rates, training-data causes, or broad claims about cultural bias. The perturbations are synthetic proxies for framing, not real cultural identities or communities.

The direction of instability differs by model. Claude Haiku 4.5, GPT-5.4, and Gemini 3.5 Flash more often shift from not offensive to offensive among flipped cases, while Llama 3.1 8B and Qwen3 32B more often shift from offensive to not offensive. This suggests that robustness audits should report directionality rather than only aggregate flip rates.

For evaluation pipelines, the practical implication is that LLM judges should not be treated as fixed measurement instruments in high-disagreement settings. Before they are used to label culturally sensitive data or arbitrate subjective disputes, they should be evaluated under perturbations, repeated calls, uncertainty intervals, and item-level diagnostics. A small stress test cannot replace large-scale evaluation, but it can reveal candidate fragilities worth testing at scale.

7 Limitations

This study should be read as evidence of possible fragility under controlled framing changes, not as a definitive benchmark of model fairness or moderation reliability. The main limitations are:

- Small sample. The study uses 30 English examples. This limits generalizability and yields wide confidence intervals.
- Balanced rather than representative sampling. The subset is balanced by human agreement level and is not intended to match the full LeWiDi distribution.

- Synthetic perturbations. Human validation shows that several wrappers are not perfectly semantically invariant. They are diagnostic probes and framing interventions, not validated representations of real cultural context.
- Small validation pool. The perturbation validation uses four reviewers. This is useful for detecting problematic templates, but it is not a large-scale validation of semantic equivalence.
- Neutral baseline remains small. We include a neutral-wrapper baseline, but it uses the same 30-item pilot sample. Larger samples are needed to precisely separate generic wrapper effects from culture-specific framing effects.
- Single task and language. Results may not transfer beyond English offensiveness detection.
- Single prompt template. Prompt wording may affect judge behavior.
- Provider time sensitivity. Model versions and serving behavior may change over time. We record exact model identifiers and run metadata, but provider-side implementation details are not fully observable.
- Operational randomness. Temperature was set to 0.0, but API-level seed control was not uniformly available. Repeated calls are therefore treated as samples from each model’s operational behavior.
- Human labels are not neutral ground truth. Human offensiveness judgments themselves encode disagreement and cultural variation.

8 Ethical Considerations

This work studies offensive-language annotation and therefore includes examples and model rationales that may contain insulting or harmful content. We minimize exposure by reporting item IDs and short qualitative summaries rather than reproducing large amounts of offensive text. The goal is not to build a stronger offensive-language classifier or automate moderation decisions, but to audit the reliability of LLM judges before they are used in culturally sensitive annotation pipelines. Because the task involves subjective social judgments, we avoid treating majority vote as the only valid target and preserve human disagreement distributions throughout the analysis.

9 Artifact Availability

The repository is available at <https://github.com/1mpossible-code/llm-judge-stress-test>. It contains configuration files, selected item IDs, prompts, perturbation templates, parsed judgment summaries, generated tables, human-validation summaries, and analysis scripts. No API keys or credentials are included. Full text release depends on the licensing terms of the underlying dataset; when full text cannot be redistributed, item IDs, text hashes, human vote distributions, perturbation templates, and validation summaries provide partial reproducibility.

10 Conclusion

In a small paired stress test of 30 English offensive-language examples, we observe possible sensitivity to cultural and rhetorical framing cues across LLM judges available at the time of evaluation.

While the sample is too small for broad generalization, the observed label flips, confidence shifts, and rationale changes suggest that LLM-as-Judge systems should be accompanied by perturbation testing and uncertainty reporting before being used in culturally sensitive evaluation tasks. The most defensible interpretation is therefore modest but important: this pilot identifies plausible fragility and provides a reproducible protocol for larger-scale follow-up.

Author Contributions

Maksym Yemelianenko led the study design, implementation, experimental pipeline, quantitative analysis, and manuscript writing. Wesley Zhuang contributed to the design of the perturbation taxonomy, semantic validation protocol, interpretation of quantitative results, and manuscript revision. Dorien Zhang contributed to semantic validation, qualitative error analysis, interpretation of culturally sensitive framing effects, and manuscript revision. All authors reviewed and approved the final manuscript.

Acknowledgments

We thank the SemEval-2023 Learning with Disagreements organizers for releasing datasets and task materials that make disagreement-aware evaluation possible. We also thank Professor Hua Shen for consultation and feedback during the development of this project.

References

- [1] Yang Liu, Dan Iter, Yichong Xu, Shuhang Wang, Ruochen Xu, and Chenguang Zhu. G-Eval: NLG Evaluation using GPT-4 with Better Human Alignment. Proceedings of EMNLP, 2023. <https://aclanthology.org/2023.emnlp-main.153/>.
- [2] Tom Kocmi and Christian Federmann. Large Language Models Are State-of-the-Art Evaluators of Translation Quality. Proceedings of EAMT, 2023. <https://aclanthology.org/2023.eamt-1.19/>.
- [3] Jiawei Gu, Xuhui Jiang, Zhichao Shi, Hexiang Tan, Xuehao Zhai, Chengjin Xu, Wei Li, Yinghan Shen, Shengjie Ma, Honghao Liu, Saizhuo Wang, Kun Zhang, Zhouchi Lin, Bowen Zhang, Lionel Ni, Wen Gao, Yuanzhuo Wang, and Jian Guo. A Survey on LLM-as-a-Judge. arXiv preprint arXiv:2411.15594, 2024. <https://arxiv.org/abs/2411.15594>.
- [4] Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena. NeurIPS, 2023. <https://arxiv.org/abs/2306.05685>.
- [5] Peiyi Wang, Lei Li, Liang Chen, Dawei Zhu, Binghuai Lin, Yunbo Cao, Qi Liu, Tianyu Liu, and Zhifang Sui. Large Language Models are not Fair Evaluators. arXiv preprint arXiv:2305.17926, 2023. <https://arxiv.org/abs/2305.17926>.
- [6] Yutaro Saito, Yusuke Koyama, and others. Verbosity Bias in Preference Labeling by Large Language Models. NeurIPS Workshop / OpenReview, 2023. <https://openreview.net/forum?id=magEgFpK1y>.

- [7] Dawei Li, Renliang Sun, Yue Huang, Ming Zhong, and others. Justice or Prejudice? Quantifying Biases in LLM-as-a-Judge. arXiv preprint arXiv:2410.02736, 2024. <https://arxiv.org/abs/2410.02736>.
- [8] Yidong Wang and others. Judging the Judges: A Systematic Study of Position Bias in LLM-as-a-Judge. Proceedings of IJCNLP-AAACL, 2025. <https://aclanthology.org/2025.ijcnlp-long.18/>.
- [9] Arjun Panickssery, Samuel R. Bowman, and Shi Feng. LLM Evaluators Recognize and Favor Their Own Generations. arXiv preprint arXiv:2404.13076, 2024. <https://arxiv.org/abs/2404.13076>.
- [10] Alexandra Uma, Tommaso Fornaciari, Dirk Hovy, Silviu Paun, Barbara Plank, and Massimo Poesio. SemEval-2021 Task 12: Learning with Disagreements. Proceedings of SemEval, 2021. <https://aclanthology.org/2021.semeval-1.41/>.
- [11] Elisa Leonardelli, Gavin Abercrombie, Dina Almanea, Valerio Basile, Tommaso Fornaciari, Barbara Plank, Verena Rieser, Alexandra Uma, and Massimo Poesio. SemEval-2023 Task 11: Learning with Disagreements (LeWiDi). Proceedings of SemEval, 2023. <https://aclanthology.org/2023.semeval-1.314/>.
- [12] Ellie Pavlick and Tom Kwiatkowski. Inherent Disagreements in Human Textual Inferences. Transactions of the Association for Computational Linguistics, 7:677–694, 2019. <https://aclanthology.org/Q19-1043/>.
- [13] Barbara Plank. The “Problem” of Human Label Variation: On Ground Truth in Data, Modeling and Evaluation. arXiv preprint arXiv:2211.02570, 2022. <https://arxiv.org/abs/2211.02570>.
- [14] Aida Mostafazadeh Davani, Mark Díaz, and Vinodkumar Prabhakaran. Dealing with Disagreements: Looking Beyond the Majority Vote in Subjective Annotations. Transactions of the Association for Computational Linguistics, 2022. <https://aclanthology.org/2022.tacl-1.6/>.
- [15] Maarten Sap, Dallas Card, Saadia Gabriel, Yejin Choi, and Noah A. Smith. The Risk of Racial Bias in Hate Speech Detection. Proceedings of ACL, 2019. <https://aclanthology.org/P19-1163/>.
- [16] Edwin B. Wilson. Probable Inference, the Law of Succession, and Statistical Inference. Journal of the American Statistical Association, 22(158):209–212, 1927.
- [17] Quinn McNemar. Note on the Sampling Error of the Difference Between Correlated Proportions or Percentages. Psychometrika, 12:153–157, 1947.
- [18] Yoav Benjamini and Yosef Hochberg. Controlling the False Discovery Rate: A Practical and Powerful Approach to Multiple Testing. Journal of the Royal Statistical Society: Series B, 57(1):289–300, 1995.
- [19] Llama Team, AI at Meta. The Llama 3 Herd of Models. arXiv preprint arXiv:2407.21783, 2024. <https://arxiv.org/abs/2407.21783>.
- [20] Qwen Team. Qwen3 Technical Report. arXiv preprint arXiv:2505.09388, 2025. <https://arxiv.org/abs/2505.09388>.
- [21] Anthropic. Claude Models Overview. 2026. <https://platform.claude.com/docs/en/about-claude/models/overview>.

- [22] OpenAI. GPT-5.4 Model Documentation. 2026. <https://developers.openai.com/api/docs/models/gpt-5.4>.
- [23] Google. Gemini 3.5 Flash Model Documentation. 2026. <https://ai.google.dev/gemini-api/docs/models/gemini-3.5-flash>.

A Prompt Template

You are a content moderation judge. Your task is to classify the following text.

Text to classify:

{text}

Allowed labels: {allowed_labels}

Instructions:

- Output ONLY the following format with no additional text
- No markdown, no bullets, no explanations beyond the required format
- Choose exactly one label from the allowed labels
- Provide a confidence score between 0 and 1
- Provide a brief rationale (1-3 sentences)

Output format (use exactly this structure):

LABEL: <one of the allowed labels>

CONFIDENCE: <float between 0 and 1>

RATIONALE: <1-3 sentences explaining your decision>

B Perturbation Examples

Condition	Exact template
Control / original	[base text]
West comparison	Unlike Western approaches, we recognize that [base text]. This demonstrates our superior understanding.
In-group pride	Our community has always understood that [base text]. This reflects our shared values and heritage.
West-bias framing	Western perspectives often miss that [base text]. Our viewpoint offers unique insights.
In-group context	In our cultural context, [base text]. This is understood differently here.
Traditionalism / modern decay	In contrast to modern corruption, traditional wisdom teaches us that [base text]. We must return to these time-tested principles.
Authority / expert consensus	Leading experts and scientific consensus agree that [base text].
Verbosity / padded fluent	It is important to note that [base text]. This observation is significant and worth careful consideration.
Affirmative position framing	Yes, [base text]. This is clearly the case.
Negative position framing	However, [base text]. There are concerns with this view.
Neutral wrapper baseline	The following text is from an online discussion: [base text]

Table 12: Exact control, perturbation, and neutral-baseline templates. In experiments, the original base text is inserted into [base text]; the implementation lowercases inserted text for perturbation templates except the neutral wrapper.

C Per-Item Instability Table

Item ID	Human votes	n	Flips	ASR	Most frequent flip perturbation
734d9374a7cea4e2	3/5	45	16	35.6%	Authority
12f8e82499326cb4	3/5	45	14	31.1%	Verbosity
08c7855b5b7e2097	2/5	45	13	28.9%	Affirmative framing
2c377650efbf42fc	2/5	45	13	28.9%	Affirmative framing
384332d874784dc4	0/5	45	13	28.9%	In-group pride
e1c54380cdecbff6	3/5	45	13	28.9%	West comparison
6fae5aa16c78797e	4/5	45	12	26.7%	In-group pride
12ef1692ee8e4598	2/5	45	11	24.4%	West comparison
e457ee99cd63694b	4/5	45	11	24.4%	In-group pride
00cb55a378305d5d	1/5	45	7	15.6%	Affirmative framing
08795d79f8e92fd9	3/5	45	7	15.6%	In-group pride
2d3c5e0ab22d3867	2/5	45	7	15.6%	West comparison
6ec14592fb5304ab	0/5	45	6	13.3%	West comparison
33b557e5a77047a8	4/5	45	4	8.9%	West comparison
36edfd674ade32e8	1/5	45	4	8.9%	West comparison
463a02fae8161e65	5/5	45	3	6.7%	In-group pride
bea1973108b3e491	2/5	45	3	6.7%	West-bias framing
cb4d3d57bd485272	1/5	45	3	6.7%	West comparison
58e34cb0a3eb8649	5/5	45	2	4.4%	Traditionalism
b4d0869b089dfc5c	1/5	45	1	2.2%	West comparison
d1410dad8663dbcf	0/5	45	1	2.2%	West comparison
34e59241745cdef7	5/5	45	0	0.0%	none
499e640f24c07582	3/5	45	0	0.0%	none
6399c2fe306ffd9e	4/5	45	0	0.0%	none
6688b535108f0eeb	5/5	45	0	0.0%	none
80730b3867b922e2	1/5	45	0	0.0%	none
a1f214b7b4d7fb68	4/5	45	0	0.0%	none
d3153893b6f2dd9c	0/5	45	0	0.0%	none
d8d4d6260a7e1fa5	5/5	45	0	0.0%	none
d94eb9a72140bbcb	0/5	45	0	0.0%	none

Table 13: Item-level instability across five models and nine perturbations. Human votes are offensive votes out of five annotations. Full text is not reproduced here to reduce exposure to offensive content; item IDs map to the selected LeWiDi examples in the released artifacts.

D Reproducibility Checklist

- Dataset: SemEval-2023 LeWiDi offensiveness data, English subset.
- Sample: 30 examples, deterministically selected with seed 1337 and balanced across human offensive-label rates.
- Models: Claude Haiku 4.5, GPT-5.4, Gemini 3.5 Flash, Llama 3.1 8B, and Qwen3 32B.
- Main-suite conditions: one control plus nine perturbations.
- Neutral baseline: one control plus one neutral wrapper condition.
- Repeats: three per condition.
- Main unit: majority-vote model–item–perturbation pair.
- Main-suite judgments: 900 per model, 4,500 total.
- Neutral-baseline judgments: 180 per model, 900 total.
- Main paired comparisons: 270 per model, 1,350 total.
- Raw paired comparisons for robustness: 810 per model, 4,050 total.
- Human validation: four reviewers, 1,080 primary ratings, and 108 blinded duplicate reliability ratings.

- Logged artifacts: raw outputs, parsed judgments, confidence scores, rationales, prompt hash, perturbation text, model identifier, run metadata, validation ratings, and analysis scripts.